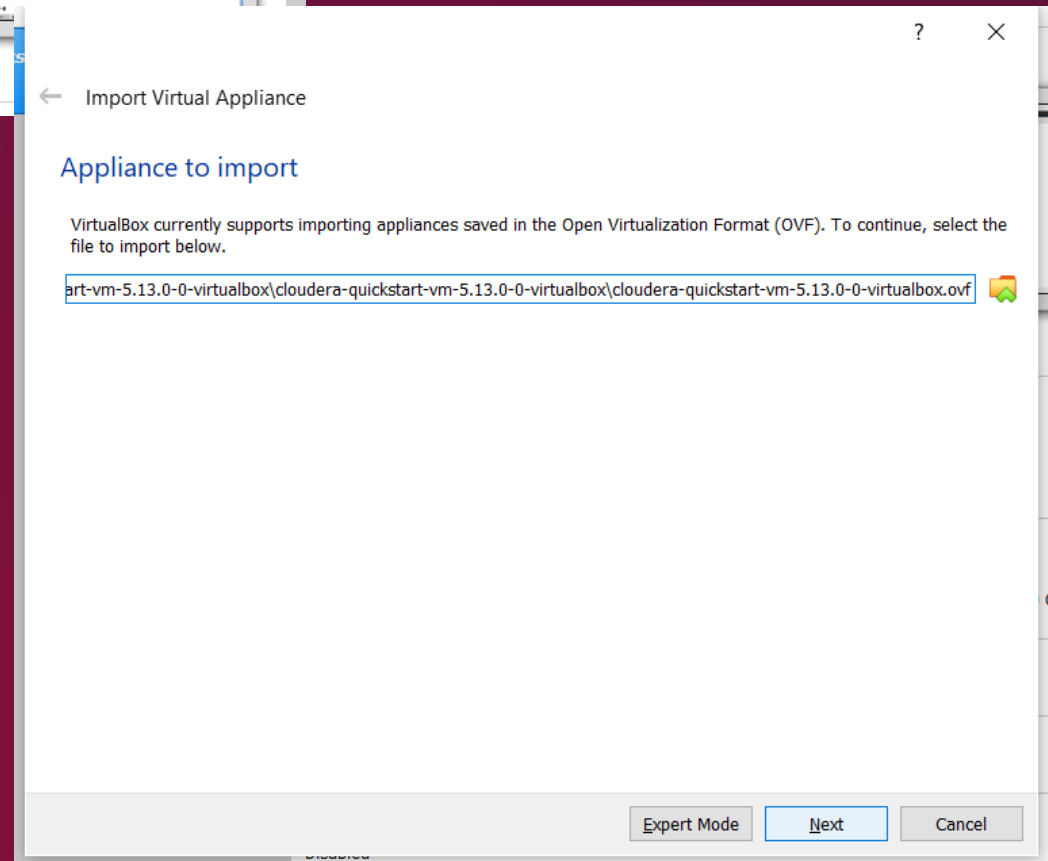
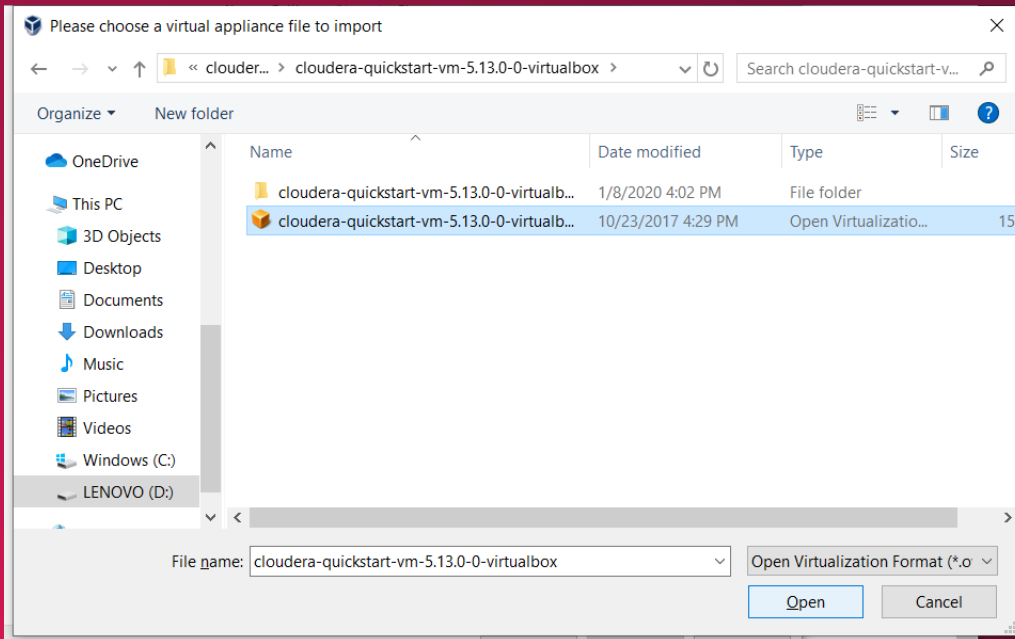
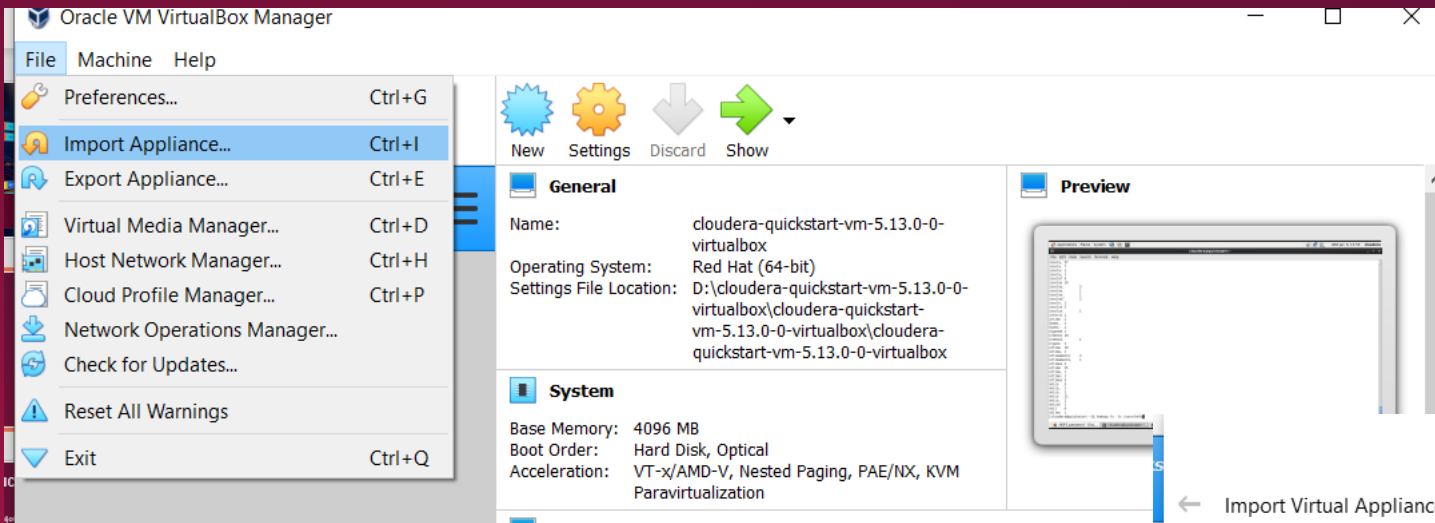



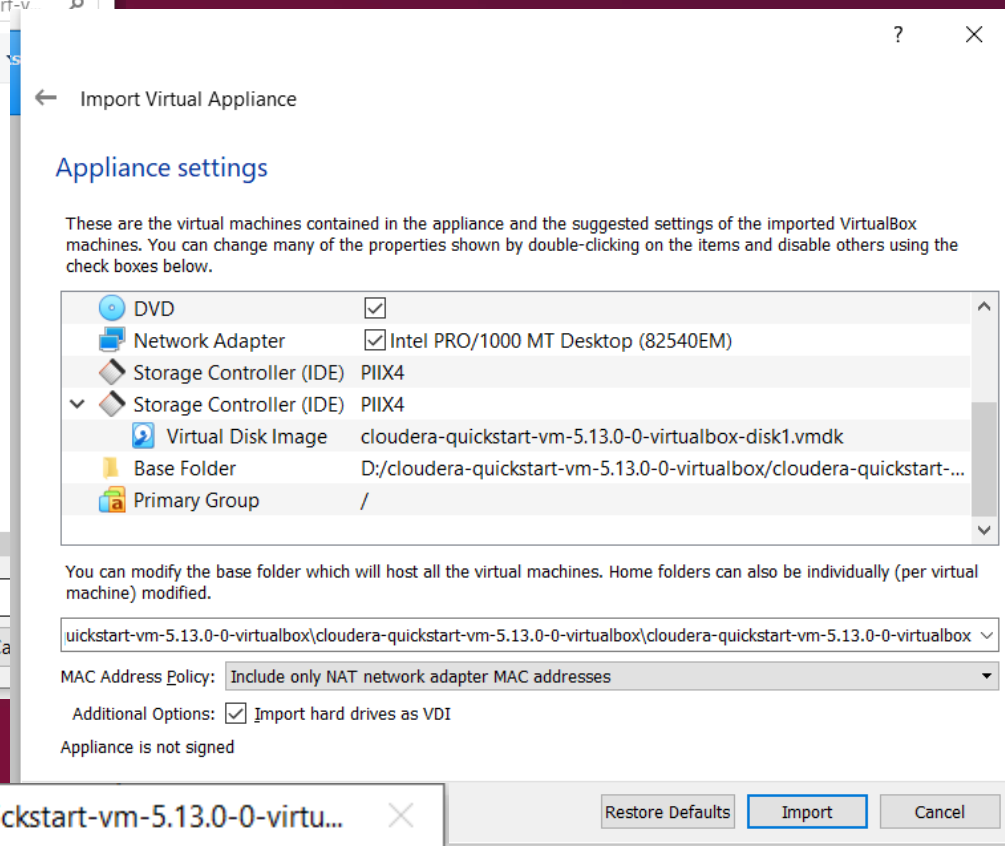
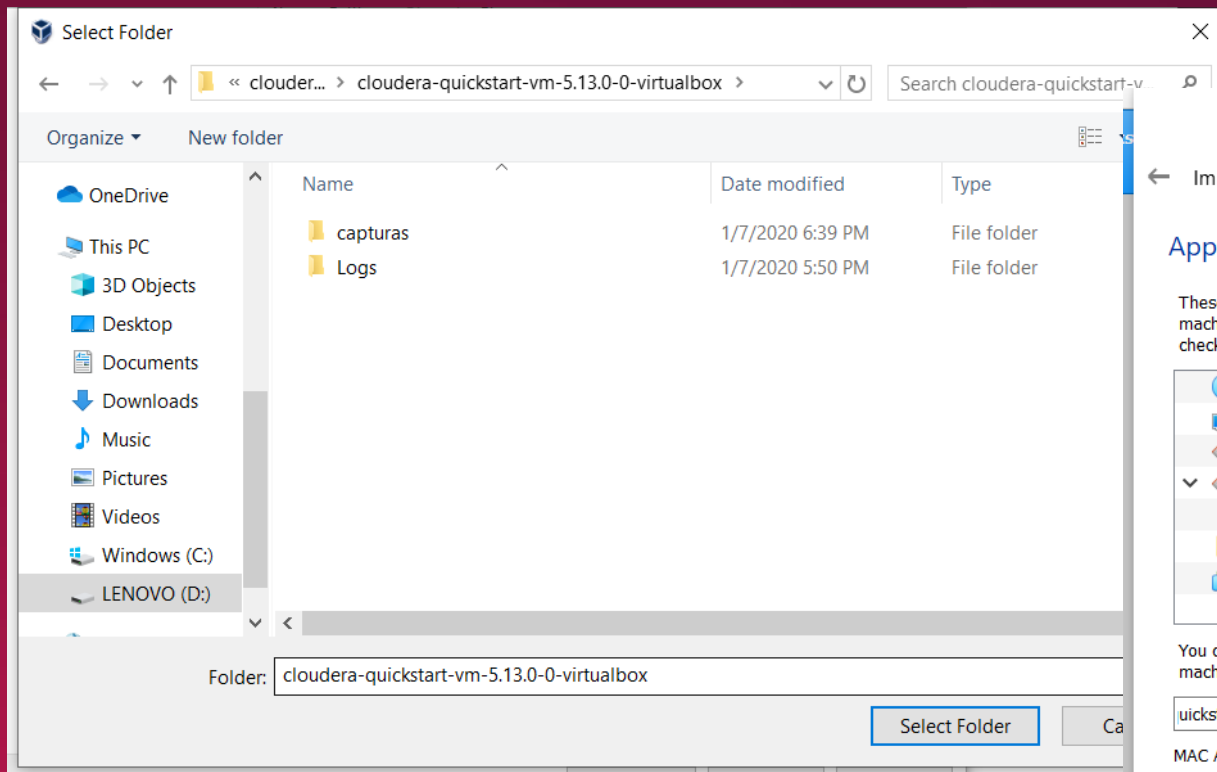


The background features a complex network of glowing blue and orange lines connecting various nodes. The nodes are labeled 'BLOCK 01', 'NODE 01', 'NODE 03', 'NODE 04', 'NODE 05', and 'BLOCK 01'. Overlaid on this network is Python code, including a class definition for 'MirrorMod' with methods 'use\_x', 'use\_y', and 'use\_z', and a function 'selection' that interacts with Blender's context and objects. The code is partially obscured by the network lines and nodes.

# Análisis de Big Data

Cindy López





Oracle VM VirtualBox Manager

FileMachineHelp

Tools

NewSettingsDiscardShow

cloudera-quickstart-vm-5.13.0-0-...  
Running

General

Name:

cloudera-quickstart-vm-5.13.0-0-virtualbox

Operating System:

Red Hat (64-bit)

Settings File Location:

D:\cloudera-quickstart-vm-5.13.0-0-virtualbox\cloudera-quickstart-vm-5.13.0-0-virtualbox\cloudera-quickstart-vm-5.13.0-0-virtualbox

System

Base Memory:

4096 MB

Boot Order:

Hard Disk, Optical

Acceleration:

VT-x/AMD-V, Nested Paging, PAE/NX, KVM Paravirtualization

Display

Video Memory:

8 MB

Graphics Controller:

VBoxVGA

Remote Desktop Server:

Disabled

Recording:

Disabled

Storage

Controller:

IDE Controller

IDE Primary Master:

cloudera-quickstart-vm-5.13.0-0-virtualbox-disk1.vdi (Normal, 64.00 GB)

IDE Secondary Master:

[Optical Drive] Empty

Audio

Disabled

Network

Adapter 1:

Intel PRO/1000 MT Desktop (NAT)

USB

Disabled

Shared folders

None

Preview

cloudera-quickstart-vm-5.13.0-0-virtualbox - Settings

GeneralSystemDisplayStorageAudioNetworkSerial PortsUSBShared FoldersUser Interface

General

BasicAdvancedDescriptionDisk Encryption

Snapshot Folder:

D:\cloudera-quickstar...-5.13.0-0-virtualbox\Snapshots

Shared Clipboard:

Bidirectional

Drag'n'Drop:

Bidirectional

Invalid settings detected

OKCancel

# Práctica MapReduce

## Contando palabras

1. Obtendremos una copia de “El Quijote” en txt.
2. Aplicaremos MapReduce.
3. Obtendremos el número de palabras que contiene nuestra copia del “El Quijote”.

# Práctica MapReduce

## Descargando nuestra copia.

1. Creamos un directorio llamado /quijote  
`cd ../../`  
`mkdir quijote`
2. Crearemos un script para descargar nuestro fichero descarga.sh  
`nano descarga.sh`  
`curl http://www.gutenberg.org/cache/epub/2000/pg2000.txt -o quijote.txt`
3. Establecemos permisos para ejecutar nuestro script  
`chmod 777 descargar.sh`
4. Ejecutamos nuestro script  
`./descarga.sh`



# Práctica MapReduce

## Descargando nuestra copia.

```
[root@quickstart /]# cd quijote/
[root@quickstart quijote]# cat descarga.sh
curl http://www.gutenberg.org/cache/epub/2000/pg2000.txt -o quijote.txt
[root@quickstart quijote]# ls -l
total 4
-rwxrwxrwx 1 root root 72 May 27 18:19 descarga.sh
[root@quickstart quijote]# ./descarga.sh
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload   Total   Spent    Left   Speed
100 2147k  100 2147k    0     0 3654k      0 --:--:-- --:--:-- --:--:-- 4121k
[root@quickstart quijote]# ls -al
total 2160
drwxr-xr-x  2 root root    4096 May 27 18:31 .
drwxrwxr-x. 23 root root    4096 May 27 18:16 ..
-rwxrwxrwx  1 root root     72 May 27 18:19 descarga.sh
-rw-r--r--  1 root root 2198927 May 27 18:31 quijote.txt
[root@quickstart quijote]#
```

# Práctica MapReduce

## Estableciendo quijote.txt en HDFS

1. `hdfs dfs -ls /user/cloudera`
2. `hdfs dfs -mkdir /user/cloudera/input`
3. `hdfs dfs -put quijote.txt /user/cloudera/input/`



# Práctica MapReduce

## Estableciendo quijote.txt en HDFS

```
[root@quickstart quijote]# hdfs dfs -ls /user/cloudera
[root@quickstart quijote]# hdfs dfs -mkdir /user/cloudera/input
[root@quickstart quijote]# hdfs dfs -put quijote.txt /user/cloudera/input/
[root@quickstart quijote]#
```

# Práctica MapReduce

Crear mapper.py –

**nano mapper.py**

```
#!/usr/bin/env python
```

```
import sys
```

```
for line in sys.stdin:
```

```
    line = line.strip()
```

```
    keys = line.split()
```

```
    for key in keys:
```

```
        value = 1
```

```
        print( "%s\t%d" % (key, value) )
```

Guardar: ctrl + o

Salir: ctrl + x

# Práctica MapReduce

## Creando nuestro mapper.py

```
[root@quickstart quijote]# cat mapper.py
#!/usr/bin/env python

import sys

for line in sys.stdin:
    line = line.strip()
    keys = line.split()
    for key in keys:
        value = 1
        print( "%s\t%d" % (key, value) )
[root@quickstart quijote]#
```

# Práctica MapReduce

Crear reducer.py – nano reducer.py

```
#!/usr/bin/env python
```

```
import sys
```

```
last_key = None
```

```
running_total = 0
```

```
for input_line in sys.stdin:  input_line =
```

```
    input_line.strip()
```

```
    this_key, value = input_line.split("\t", 1)
```

```
    value = int(value)
```

# Práctica MapReduce

## Creando nuestro reducer.py

```
    if last_key == this_key:
        running_total += value
    else:
        if last_key:
            print( "%s\t%d" % (last_key, running_total) )
            running_total = value
        last_key = this_key
```

```
if last_key == this_key:
    print( "%s\t%d" % (last_key, running_total) )
```

Guardar: ctrl + o

Salir: ctrl + x

# Práctica MapReduce

## Creando nuestro reducer.py

```
[root@quickstart quijote]# cat reducer.py
#!/usr/bin/env python

import sys

last_key = None
running_total = 0

for input_line in sys.stdin:
    input_line = input_line.strip()
    this_key, value = input_line.split("\t", 1)
    value = int(value)

    if last_key == this_key:
        running_total += value
    else:
        if last_key:
            print( "%s\t%d" % (last_key, running_total) )
            running_total = value
            last_key = this_key

if last_key == this_key:
    print( "%s\t%d" % (last_key, running_total) )
[root@quickstart quijote]#
```



# Práctica MapReduce

## Ejecutando nuestros archivos

1. `chmod 777 *.py`
2. `hdfs dfs -mkdir /user/cloudera/input`
3. `hdfs dfs -put quijote.txt /user/cloudera/input/`
4. `hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar -  
input /user/cloudera/input -output /user/cloudera/output -  
mapper /quijote/mapper.py -reducer /quijote/reducer.py`

# Práctica MapReduce

## Ejecutando nuestros archivos

```
[root@quickstart quijote]# hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming-2.6.0-cdh5.13.0.jar -input /user/cloudera/input -output /user/cloudera/output6 -mapper /quijote/mapper.py -reducer /quijote/reducer.py
packageJobJar: [] [/usr/lib/hadoop-mapreduce/hadoop-streaming-2.6.0-cdh5.13.0.jar] /tmp/streamjob4301095940548298469.jar tmpDir=null
18/05/28 11:23:28 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032
18/05/28 11:23:28 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032
18/05/28 11:23:29 INFO mapred.FileInputFormat: Total input paths to process : 1
18/05/28 11:23:29 INFO mapreduce.JobSubmitter: number of splits:2
18/05/28 11:23:29 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1527380325361_0007
18/05/28 11:23:29 INFO impl.YarnClientImpl: Submitted application application_1527380325361_0007
18/05/28 11:23:29 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1527380325361_0007/
18/05/28 11:23:29 INFO mapreduce.Job: Running job: job_1527380325361_0007
18/05/28 11:23:37 INFO mapreduce.Job: Job job_1527380325361_0007 running in uber mode : false
18/05/28 11:23:37 INFO mapreduce.Job:  map 0% reduce 0%
18/05/28 11:23:49 INFO mapreduce.Job:  map 50% reduce 0%
18/05/28 11:23:50 INFO mapreduce.Job:  map 100% reduce 0%
```

# Práctica MapReduce

## Visualizando resultados

### 1. `hdfs dfs -ls /user/cloudera/output`

```
[root@quickstart quijote]# hdfs dfs -ls /user/cloudera/output6
Found 2 items
-rw-r--r--  1 root cloudera          0 2018-05-28 11:23 /user/cloudera/output6/
SUCCESS
-rw-r--r--  1 root cloudera    448894 2018-05-28 11:23 /user/cloudera/output6/
part-00000
[root@quickstart quijote]#
```

# Práctica MapReduce

## Visualizando resultados

1. `hdfs dfs -cat /user/cloudera/output/part-00000 | head -1000`
2. `hdfs dfs -cat /user/cloudera/output/*`

```
[root@quickstart quijote]# hdfs dfs -cat /user/cloudera/output6/part-00000 | head -1000
!Mal      1
"Al       1
"Cuando  2
"Cuidados      1
"De       2
"Defects,"    1
"Desnudo     1
"Dijo       1
"Dime       1
"Don        1
"Donde      1
"Dulcinea    1
"El         2
"Esta       1
"Harto      1
"Iglesia,    1
"Information 1
"Más        2
"No         5
"Nunca      1
"Plain      2
"Project    5
```