

Frankfurt University of Applied Sciences
Prof. Dr. Christina Andersson

Solutions to R Exercises

1 Introduction

2 Manipulation of Vectors and Numbers

1. (a) `> x=c(8,9,7,5)`
(b) `> y=c(1,2,3,4)`
(c) `> s=x+y`
(d) `> s[2]`
[1] 11
2. (a) `> p=matrix(c(1,2,3,3,5,1,4,2), nrow=4, ncol=2,byrow=T)`
(b) `> p`
[,1] [,2]
[1,] 1 2
[2,] 3 3
[3,] 5 1
[4,] 4 2
(c) `> p[,2]`
[1] 2 3 1 2
(d) `> p[3,2]`
[1] 1
3. (a) `> x=c(5,6,7)`
(b) `> max(x)`
[1] 7
(c) `> rev(x)`
[1] 7 6 5
4. (a) `> s=seq(6,76, by=1)`
`> s`
[1] 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30
[26] 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55
[51] 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76
(b) `> sample(s, 5, replace = F)`
[1] 66 17 51 59 32
(c) `> sample(s, 5, replace = T)`
[1] 63 47 15 15 45
5. (a) `> d=seq(0,99,by=1)`
(b) `> e=rnorm(100,3,4)`
(c) `> f = d + e`
6. `> x = c(2,3,4)`
`> y = c(5,6,7)`
`> test = data.frame(x,y)`
`> test`
`> test`
x y

```

1 2 5
2 3 6
3 4 7

```

```

7. (a) > library(datasets)
      > nrow(airquality)
      [1] 153
      > ncol(airquality)
      [1] 6

```

```

(b) > head(airquality)
      Ozone Solar.R Wind Temp Month Day
1      41      190  7.4   67     5   1
2      36      118  8.0   72     5   2
3      12      149 12.6   74     5   3
4      18      313 11.5   62     5   4
5      NA       NA 14.3   56     5   5
6      28       NA 14.9   66     5   6

```

3 Tables and Graphs

```

1. (a) exam = c(0,0,0,0,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,0,0,1,0,0,0,0,0,0,0,
1,0,0,2,0,0,0,0,0,0,0,0,1,0,0,0,1,0,0,0,0,0,1,0,0,1,0,0,3,1,1,2,0,0,1,
0,0,0,0,2,0,1,0,1,0,1,0,0,0,0,0,1,0,0,1,0,0,1,1,0,1,0,0,1,0,0,0,0,0,1,
0,1,1,1,2,0,0,1,0,2,1,1,0,0,0,0,0,1,0,1,1,0,0,0,2,1,1,1,0,2,0,1,1,2,
1,1,0,1,0)

```

```

abs_freq = table(exam)
> abs_freq
exam
 0  1  2  3
90 37  8  1

```

```


(b) > rel_freq = abs_freq/length(exam)
> rel_freq
exam
      0      1      2      3
0.661764706 0.272058824 0.058823529 0.007352941

```


```

(c) > barplot(abs_freq,names.arg=c("0","1","2","3"),main="Bar plot",
      xlab="No of trials",ylab="Frequencies")

```



```
(d) fn = cumsum(rel_freq)
plot(sort(unique(exam)),fn,type="n",xlab="Trials",
ylab="Fn",main="Distribution Function")
lines(sort(unique(exam)),fn,type="s")
```




2. (a)

```
> umbrella = data.frame(no_of_umbrellas=c(0,3,4,5,6,8),days=c(4,6,8,3,1,2))
> abs_freq = umbrella$days
> abs_freq
[1] 4 6 8 3 1 2
> cum_abs = cumsum(abs_freq)
> cum_abs
[1] 4 10 18 21 22 24
```
- (b)

```
> rel_freq = umbrella$days/sum(abs_freq)
> rel_freq
[1] 0.1666667 0.2500000 0.3333333 0.1250000 0.0416667 0.0833333
> cum_rel <- cumsum(rel_freq)
> cum_rel
[1] 0.1666667 0.4166667 0.7500000 0.8750000 0.9166667 1.0000000
```
- (c)

```
> fn = cumsum(rel_freq)
> plot(umbrella$no_of_umbrellas,fn,type="n",xlab="Number of umbrellas",
+ ylab="Fn",main="Distribution Function")
> lines(umbrella$no_of_umbrellas,fn,type="s")
```




(d) `> fn[5]`

`[1] 0.9166667`


During 91.67 % of the days 6 umbrellas or less are sold.

3. (a) `boots = data.frame(Year=1985:2008, No_of_Accidents=c(472787, 495871, 532220, 523387, 499666, 513587, 487654, 478721, 521972, 476544, 430432, 452165, 432589, 456436, 457422, 466064, 519482, 343091, 328221, 522169, 415077, 387212, 415254, 423731))`

(b) `> barplot(table(boots$No_of_Accidents))`




- (c) `hist(boots$No_of_Accidents, breaks=seq(min(boots$No_of_Accidents), max(boots$No_of_Accidents), (max(boots$No_of_Accidents)-min(boots$No_of_Accidents))/9), freq=F, main="Histogram", xlab="Number of Accidents", col="grey")`




In the last class.


```
4. > x=c(1,2,2,3,3,4,4,5,6,6,7,7,8,9,8,9,9)
> y =c(1.426865, 2.495512, 2.751945, 3.794935, 3.682121, 3.692246,
4.451148, 5.200307, 5.638318, 7.672076, 6.819001, 7.208195, 9.076866,
9.441328, 7.752522, 9.545205, 10.097847)
> plot(x,y)
```




```
5. > group1 = c(12, 23, 34, 33, 35, 33, 32, 31, 30, 29, 28, 28, 27, 27, 6, 26)
> group2 = c(13, 13, 24, 30, 31, 31, 30, 30, 31, 28, 28, 29, 26, 25, 26, 26)
> boxplot(group1, group2)
```



6. (a) `Gr1=c(22,22,28,23,45,21,22)`
`Gr2=c(21,23,21,24,22,20,21)`
`Gr3=c(48,45,51,29,38,40,38)`
(b) `boxplot(Gr1, Gr2, Gr3)`
Interpretation: -




7. (a) `> library(datasets)`
(b) `> boxplot(beaver1$temp, beaver2$temp)`




Interpretation?

8. (a) `> library(MASS)`
(b) `> max(iris$Sepal.Width)`
`[1] 4.4`
(c) `> plot(iris$Sepal.Length, iris$Sepal.Width)`



(d) `> plot(iris$Sepal.Length,iris$Sepal.Width,col=iris$Species)`



4 Measures of Central Tendency

```
1. > strike= data.frame(Country=c("Aland", "Bland", "Cland", "Dland"),
  Days=c(77, 45, 76, 83))
```

```
> strike
```

	Country	Days
1	Aland	77
2	Bland	45
3	Cland	76
4	Dland	83

```
(a) > median(strike$Days)
[1] 76.5
```

```
(b) > mean(strike$Days)
[1] 70.25
```

```
(c) -
```

```
2. > students=data.frame(name=c("Anton", "Kim", "Harald", "Inga", "Mona", "Sigrid"),
+ height=c(170,167,169,172,171,170), weight=c(70,75,120,87,88,87))
```

```
> students
```

	name	height	weight
1	Anton	170	70
2	Kim	167	75
3	Harald	169	120
4	Inga	172	87
5	Mona	171	88
6	Sigrid	170	87


```

1 Anton    170    70
2 Kim      167    75
3 Harald   169   120
4 Inga     172    87
5 Mona     171    88
6 Sigrid   170    87

```

```

(a) > mean(students$weight)
[1] [1] 87.83333
> median(students$weight)
[1] 87

```

```

(b) > students2=subset(students,name!="Harald")
> students2
      name height weight
1 Anton    170    70
2 Kim      167    75
4 Inga     172    87
5 Mona     171    88
6 Sigrid   170    87
> mean(students2$weight)
[1] 81.4
> median(students2$weight)
[1] 87

```

(c) -

3. `grades<-c(2,3,3,3,4,1,5,2,4,2,2,2,3,4,4,3,2,1,3,3,3,2,2)`

```

(a) > table(grades)
grades
1 2 3 4 5
2 8 8 4 1

```

```

(b) > mean(grades)
[1] 2.73913
> median(grades)
[1] 3

```

From the table, we see that the modes are 2 and 3.

```

4. > pumpkins=data.frame(no_of_pumpkins=c(0,1,2,3,4,5,6,7),
days=c(10,2,3,5,4,2,4,5))
> sum(pumpkins$no_of_pumpkins*pumpkins$days)/sum(pumpkins$days)
[1] 3.085714

```

5. (a) `> library(datasets)`

```

(b) > mean(airquality$Temp)
[1] 77.88235

```

```

(c) > median(airquality$Solar.R)
[1] NA

```

```

(d) > median(airquality$Solar.R, na.rm=T)
[1] 205

```

```
6. > (90*0+4*1+6*2)/100
[1] 0.16
```

The average number of server problems per day is 0.16.

```
7. > flats=data.frame(rooms=1:7,no_of_flats=c(56,55,35,22,12,6,2))
> sum(flats$rooms*flats$no_of_flats)/sum(flats$no_of_flats)
[1] 2.494681
```

```
8. > mydata=data.frame(obs_values=c(50,45,25,20),frequency=c(2,4,2,2))
> prod(mydata$obs_values[1]^mydata$frequency[1],
mydata$obs_values[2]^mydata$frequency[2],
mydata$obs_values[3]^mydata$frequency[3],
mydata$obs_values[4]^mydata$frequency[4])^(1/sum(mydata$frequency))
[1] 34.74346
```

```
9. > mydata=c(5,7,8,9,9)
> prod(mydata)^(1/length(mydata))
[1] 7.432392
```

```
10. > mydata=c(5,7,8,9,9)
> length(mydata)/sum(1/mydata)
> [1] 7.245543
```

```
11. > mydata = c(5,8,9,9,9,4,5,6,6,76,43,56,65,65,3,34,45)
> quantile(mydata)
 0%  25%  50%  75% 100%
  3    6    9   45   76
```

```
12. > scores = c(43,12,11,22,23,34,34,33,34,23,33,32,11,9,45,
56,48,23,23,43,23,21,21,45,23,22,32,32,21,43,11,47)
> quantile(scores, 0.96)
 96%
47.76
```

Those students that had scores 56 and 48.

5 Measures of Spread

```
1. (a) > scores = c(56, 87, 88, 91, 66)
> sd(scores)
[1] 15.6301
```

```
(b) > var(scores)
[1] 244.3
```

```
(c) > range(scores)
[1] 56 91
> 91-56
[1] 35
```

```
2. (a) > library(datasets)
(b) > var(beaver1$temp)
[1] 0.03741196
> var(beaver2$temp)
[1] 0.1996203
```

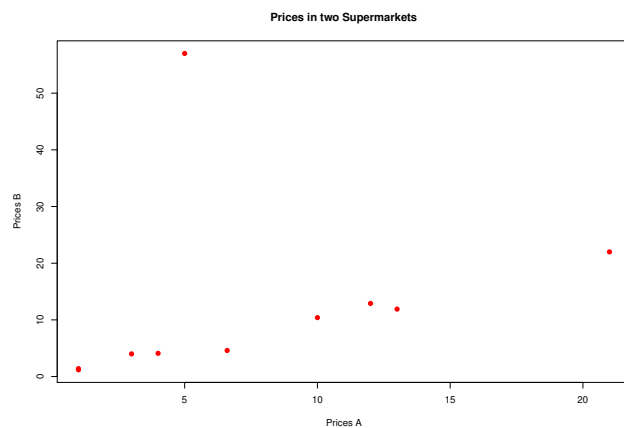
(c) -

```
3. > datawith=c(17, 23, 33, 24, 78)
> datawithout=c(17, 23, 33, 24)
> iqr_with=quantile(datawith, 0.75)-quantile(datawith, 0.25)
> iqr_with
75%
10
> iqr_without=quantile(datawithout, 0.75)-quantile(datawithout, 0.25)
> iqr_without
75%
4.75
> range_with=max(datawith)-min(datawith)
> range_with
[1] 61
> range_without=max(datawithout)-min(datawithout)
> range_without
[1] 16
```

6 Correlation

```
1. (a) > A=c(1, 5, 6.6, 4, 10, 12, 13, 21, 1, 3)
> B=c(1.2, 57, 4.6, 4.1, 10.4, 12.9, 11.9, 22, 1.4, 4)
> cor(A,B)
[1] 0.2399742
```

(b) Scatter plot:

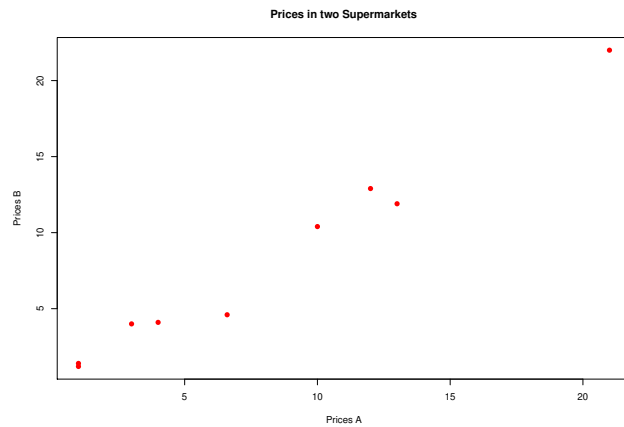


```
> plot(A,B,pch=16,cex=1.2,col="red",main="Prices in two Supermarkets",
xlab="Prices A",ylab="Prices B")
```

(c) -

```
(d) > Anew=c(1, 6.6, 4, 10, 12, 13, 21, 1, 3)
> Bnew=c(1.2, 4.6, 4.1, 10.4, 12.9, 11.9, 22, 1.4, 4)
> cor(Anew,Bnew)
[1] 0.9889279
```

Scatter plot:



```
> plot(Anew,Bnew,pch=16,cex=1.2,col="red",main="Prices in two Supermarkets",
      xlab="Prices A",ylab="Prices B")
```

2. (a)

```
> y=c(4,4,7,11,11)
```




```
> x=c(1,2,3,5,10)
```



```
> plot(x,y,xlab="x",ylab="y")
```


 Scatter plot:



(b)

```
> cor(x,y)
```



```
[1] 0.8521091
```



```
> cor(x,y, method="spearman")
```



```
[1] 0.9486833
```

(c)

```
> rank(x)
```



```
[1] 1 2 3 4 5
```



```
> rank(y)
```



```
[1] 1.5 1.5 3.0 4.5 4.5
```

Double values exist.

3.

```
> W=c(1,2,3,3,6,1,3,8)
```



```
> V=c(4,4,5,3,8,1,5,6)
```



```
> cor(W,V, method="spearman")
```



```
[1] 0.8013814
```

7 Regression

```
1. > d=c(1,1,1,2,2,2,3,3,5,6,7,8)
   > e=c(2,3,4,4,5,6,6,7,8,8,8,9)
```

```
(a) > lm(e ~ d)
```

```
Call:
```

```
lm(formula = e ~ d)
```

```
Coefficients:
```

```
(Intercept)          d
      3.0336      0.8194
```

```
i.e.  $e = 3.0336 + 0.8194 * d$ 
```

```
(b) ...
```

```
(c) > cor(d,e)
     > 0.898421
```

```
(d) ...
```

```
(e) > cov(d,e)
     > 4.984848
```

```
(f) ...
```

```
2. (a) > cor(d,e, method='spearman')
     [1] -0.02683367
```

```
(b) > cor(d,e)
     [1] 0.1627058
```

```
(c) ...
```

```
(d) ...
```

```
3. (a) e=c(5,2,3,4,2,1,5,4,7,5,8,9,8,8)
     j=c(5,6,3,4,1,1,1,6,7,8,8,7,8,9)
```

```
> lm(e~j)
```

```
Call:
```

```
lm(formula = e ~ j)
```

```
Coefficients:
```

```
(Intercept)          j
      1.5068      0.6744
```

```
> summary(lm(e~j))
```

```
Call:
```

```
lm(formula = e ~ j)
```

```
Residuals:
```

```
      Min       1Q   Median       3Q      Max
-3.553 -1.018 -0.030  1.017  2.819
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.5068	1.0511	1.434	0.17724
j	0.6744	0.1766	3.819	0.00244 **

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 1.808 on 12 degrees of freedom


Multiple R-squared: 0.5486, Adjusted R-squared: 0.511

F-statistic: 14.58 on 1 and 12 DF, p-value: 0.002444

```
> plot(residuals(lm(e~j)))
```

The estimated regression line is $e = 1.5068 + 0.6744 * j$. The p-value 0.00244 tells that the variable j is significant in the model. The value of the multiple R-squared tells that the model explains 54.86% of the variation in e .

The residual plot looks rather random.



(b) -

(c) Take the square root of 0.5486 to get the correlation.

8 Probability Distributions

1. (a) `> pbinom(35, 100, 0.4)`

```
[1] 0.1794694
```

(b) $P(X > 39) = 1 - P(X \leq 39)$

```
> 1-pbinom(39, 100, 0.4)
```

```
[1] 0.5379247
```

(c) `> sum(dbinom(36:38, size = 100, prob = 0.4))`

```
[1] 0.2027183
```

2. (a) $P(X \leq 11)$

```
> pnorm(11, mean = 10, sd = 4)
```

```
[1] 0.5987063
```

(b) $P(X > 13)$

```

> 1 - pnorm(13, mean = 10, sd = 4)
[1] 0.2266274
(c)  $P(10 \leq X \leq 12)$ 
> pnorm(12, mean = 10, sd = 4) - pnorm(10, mean = 10, sd = 4)
[1] 0.1914625
3. > qnorm(0.50, mean=0, sd=10)
[1] 0
4. > rexp(10, 5)
[1] 0.091088054 0.153399162 0.444140176 0.075040071 0.136310121 0.057523302
[7] 0.008389491 0.089599280 0.046229349 0.183981225
5. > qbinom(0.4, 100, 0.3)
[1] 29
6. > 1 - pnorm(36, mean=35.42, sd=sqrt(16))
[1] 0.4423554

```

9 Hypothesis Tests

1. We summarize the values needed for the test statistic:

```

> xbar = 14.6
> mu0 = 15.4
> sigma = 2.5
> n = 35
> z =(xbar - mu0)/(sigma/sqrt(n))
> z
[1] -1.893146

```

Then, we compute the critical values using *qnorm*:

```

> qnorm(0.05/2)
[1] -1.959964

```

this means that the critical values are aprox. -1.96 and 1.96.

Since the value of the test statistic is $-1.96 < -1.89 < 1.96$, this means that we cannot reject the null hypothesis at significance level 5%. We cannot reject that the mean weight is 15.4 this year.

2. > 2* pnorm(-1.893146)


```
[1] 0.05833846
```

where we used the value for $z = -1.893146$ from previous exercise. We cannot reject the null hypothesis, since the p -value is larger than 5%.

3. Hypotheses:

$$\begin{aligned}
 H_0: & \mu \geq 10,000 \\
 H_1: & \mu < 10,000
 \end{aligned}$$

We summarize the values needed for the test statistic:

```
> xbar = 9900
> mu0 = 10000
> sigma = 125
> n = 30
> z =(xbar - mu0)/(sigma/sqrt(n))
> z
[1] -4.38178
```

Then, we compute the critical values using *qnorm*:

```
> qnorm(0.05)
[1] -1.644854
```

This means that the critical value is aprox. -1.64. Since $-4.38178 < -1.64$, we can reject the null hypothesis at level 5%.

4. Using the value from previous exercise, a lower tail p -value is:

```
> pnorm(-4.38178)
[1] 5.885682e-06
```

which means that we can reject the null hypothesis.

5. (a) `> library(datasets)`
(b) `> mean(beaver1$temp)`
[1] 36.86219
(c) `> t.test(beaver1$temp, mu=37)`
One Sample t-test

```
data: beaver1$temp
t = -7.6071, df = 113, p-value = 9.038e-12
alternative hypothesis: true mean is not equal to 37
95 percent confidence interval:
 36.82630 36.89808
sample estimates:
mean of x
 36.86219
```

The p -value tells that we can reject the null hypothesis.

6. `> x = matrix(c(208, 230, 282, 241), ncol = 2)`
`> x`
[,1] [,2]
[1,] 208 282
[2,] 230 241
`> chisq.test(x)`

Pearson's Chi-squared test with Yates' continuity correction

```
data: x
```


X-squared = 3.6919, df = 1, p-value = 0.05468

Since the p -value $0.05468 > 0.05$, we cannot reject the null hypothesis. We cannot reject that the two variables can be independent.

```
7. > library(MASS)
> freq_data<-table(Aids2$state, Aids2$status)
> freq_data
```

	A	D
NSW	664	1116
Other	107	142
QLD	78	148
VIC	233	355

```
> chisq.test(freq_data)
```

Pearson's Chi-squared test

data: freq_data

X-squared = 4.7982, df = 3, p-value = 0.1872

Since the p -value $0.1872 > 0.05$, we cannot reject the null hypothesis. We cannot reject that the two variables can be independent.

10 Confidence Intervals

```
1. (a) > library(datasets)
(b) > mean(beaver1$temp)
[1] 36.86219
(c) > t.test(beaver1$temp, conf.level=0.99)
One Sample t-test
```

data: beaver1\$temp

t = 2034.8, df = 113, p-value < 2.2e-16

alternative hypothesis: true mean is not equal to 0

99 percent confidence interval:

36.81473 36.90966

sample estimates:

mean of x

36.86219

The confidence interval is [36.81473, 36.90966]